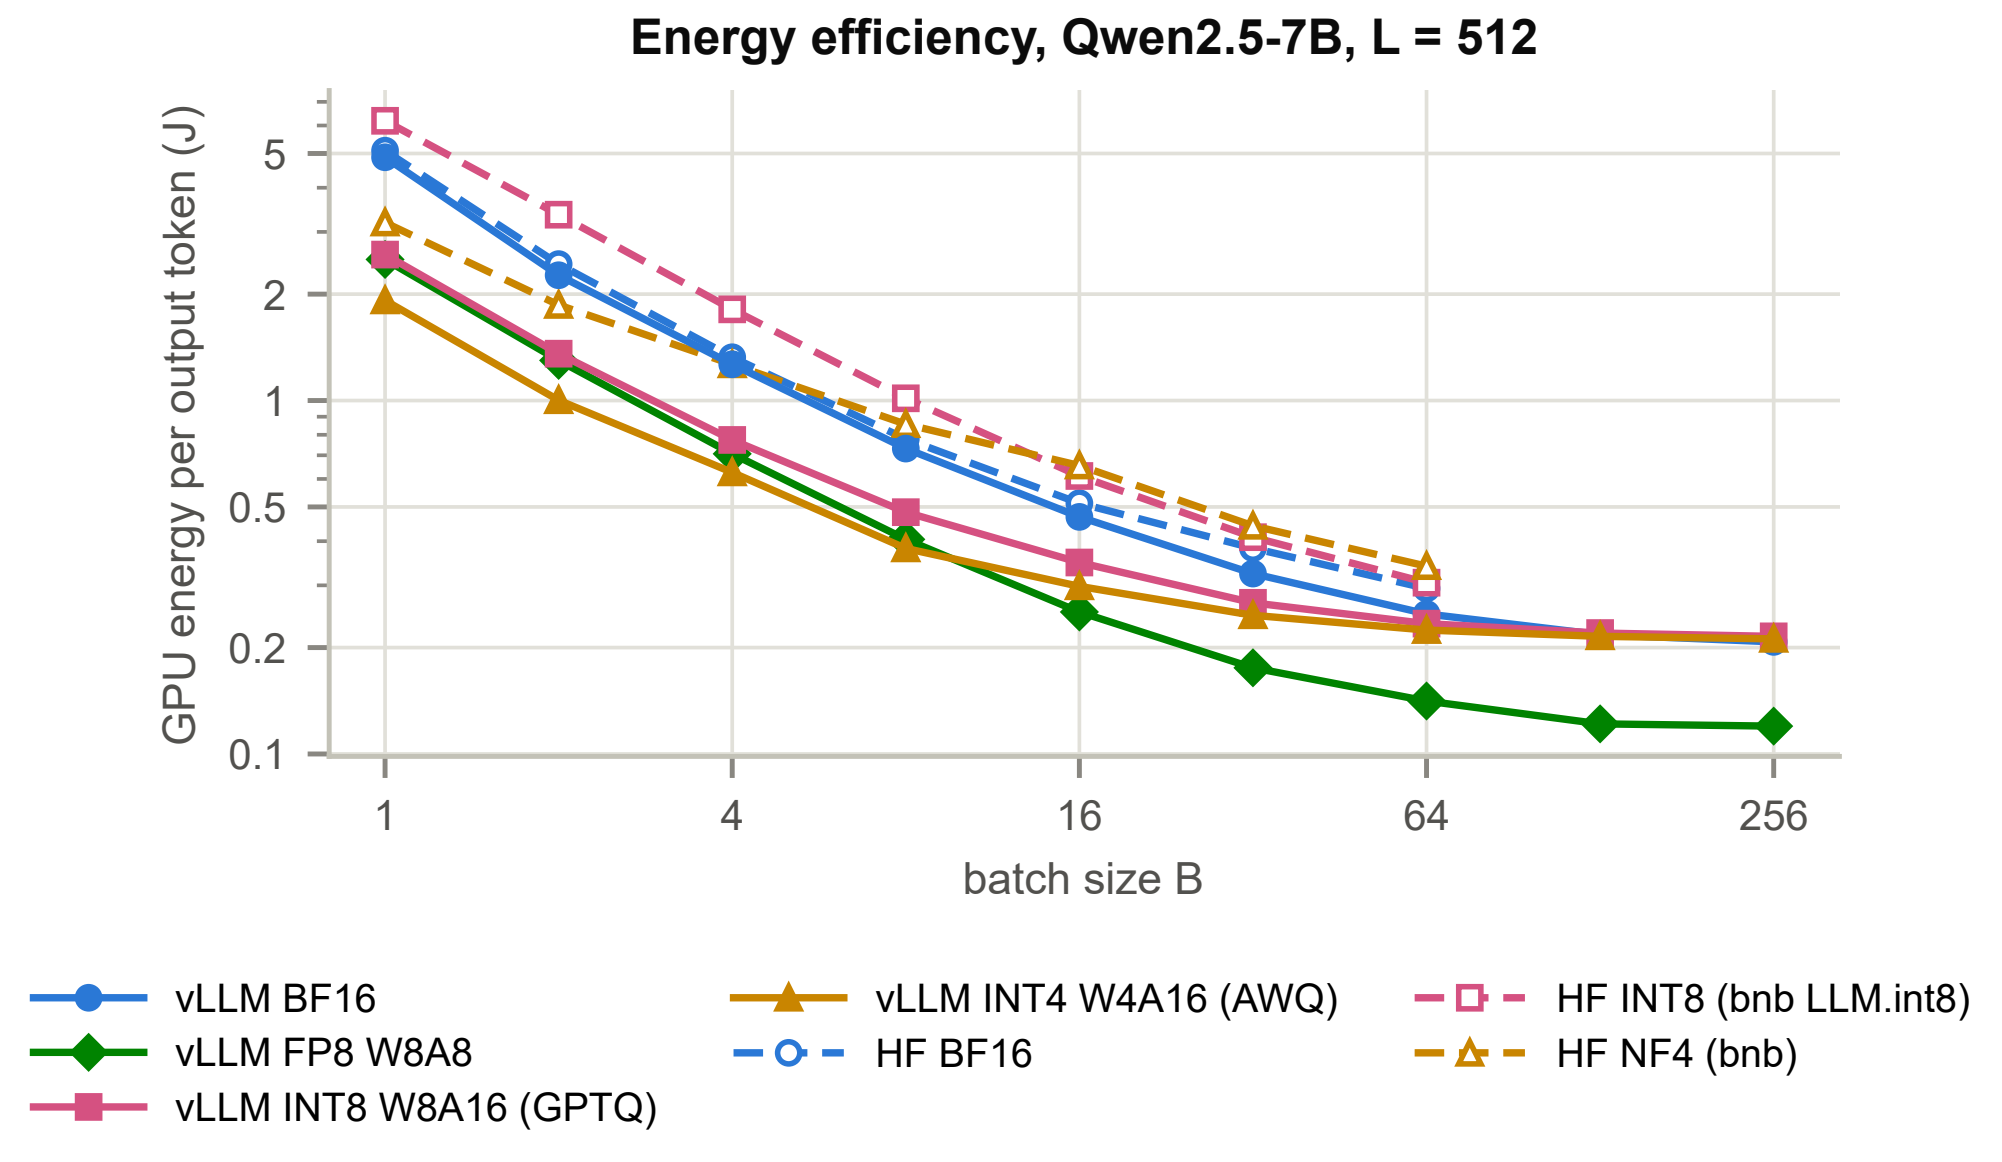
Energy efficiency, Qwen2.5-7B, L = 512
GPU energy per output token (J)
5
2
1
0.5
0.2
0.1
1
4
16
64
256
batch size B
vLLM BF16
vLLM FP8 W8A8
vLLM INT8 W8A16 (GPTQ)
vLLM INT4 W4A16 (AWQ)
HF BF16
HF INT8 (bnb LLM.int8)
HF NF4 (bnb)