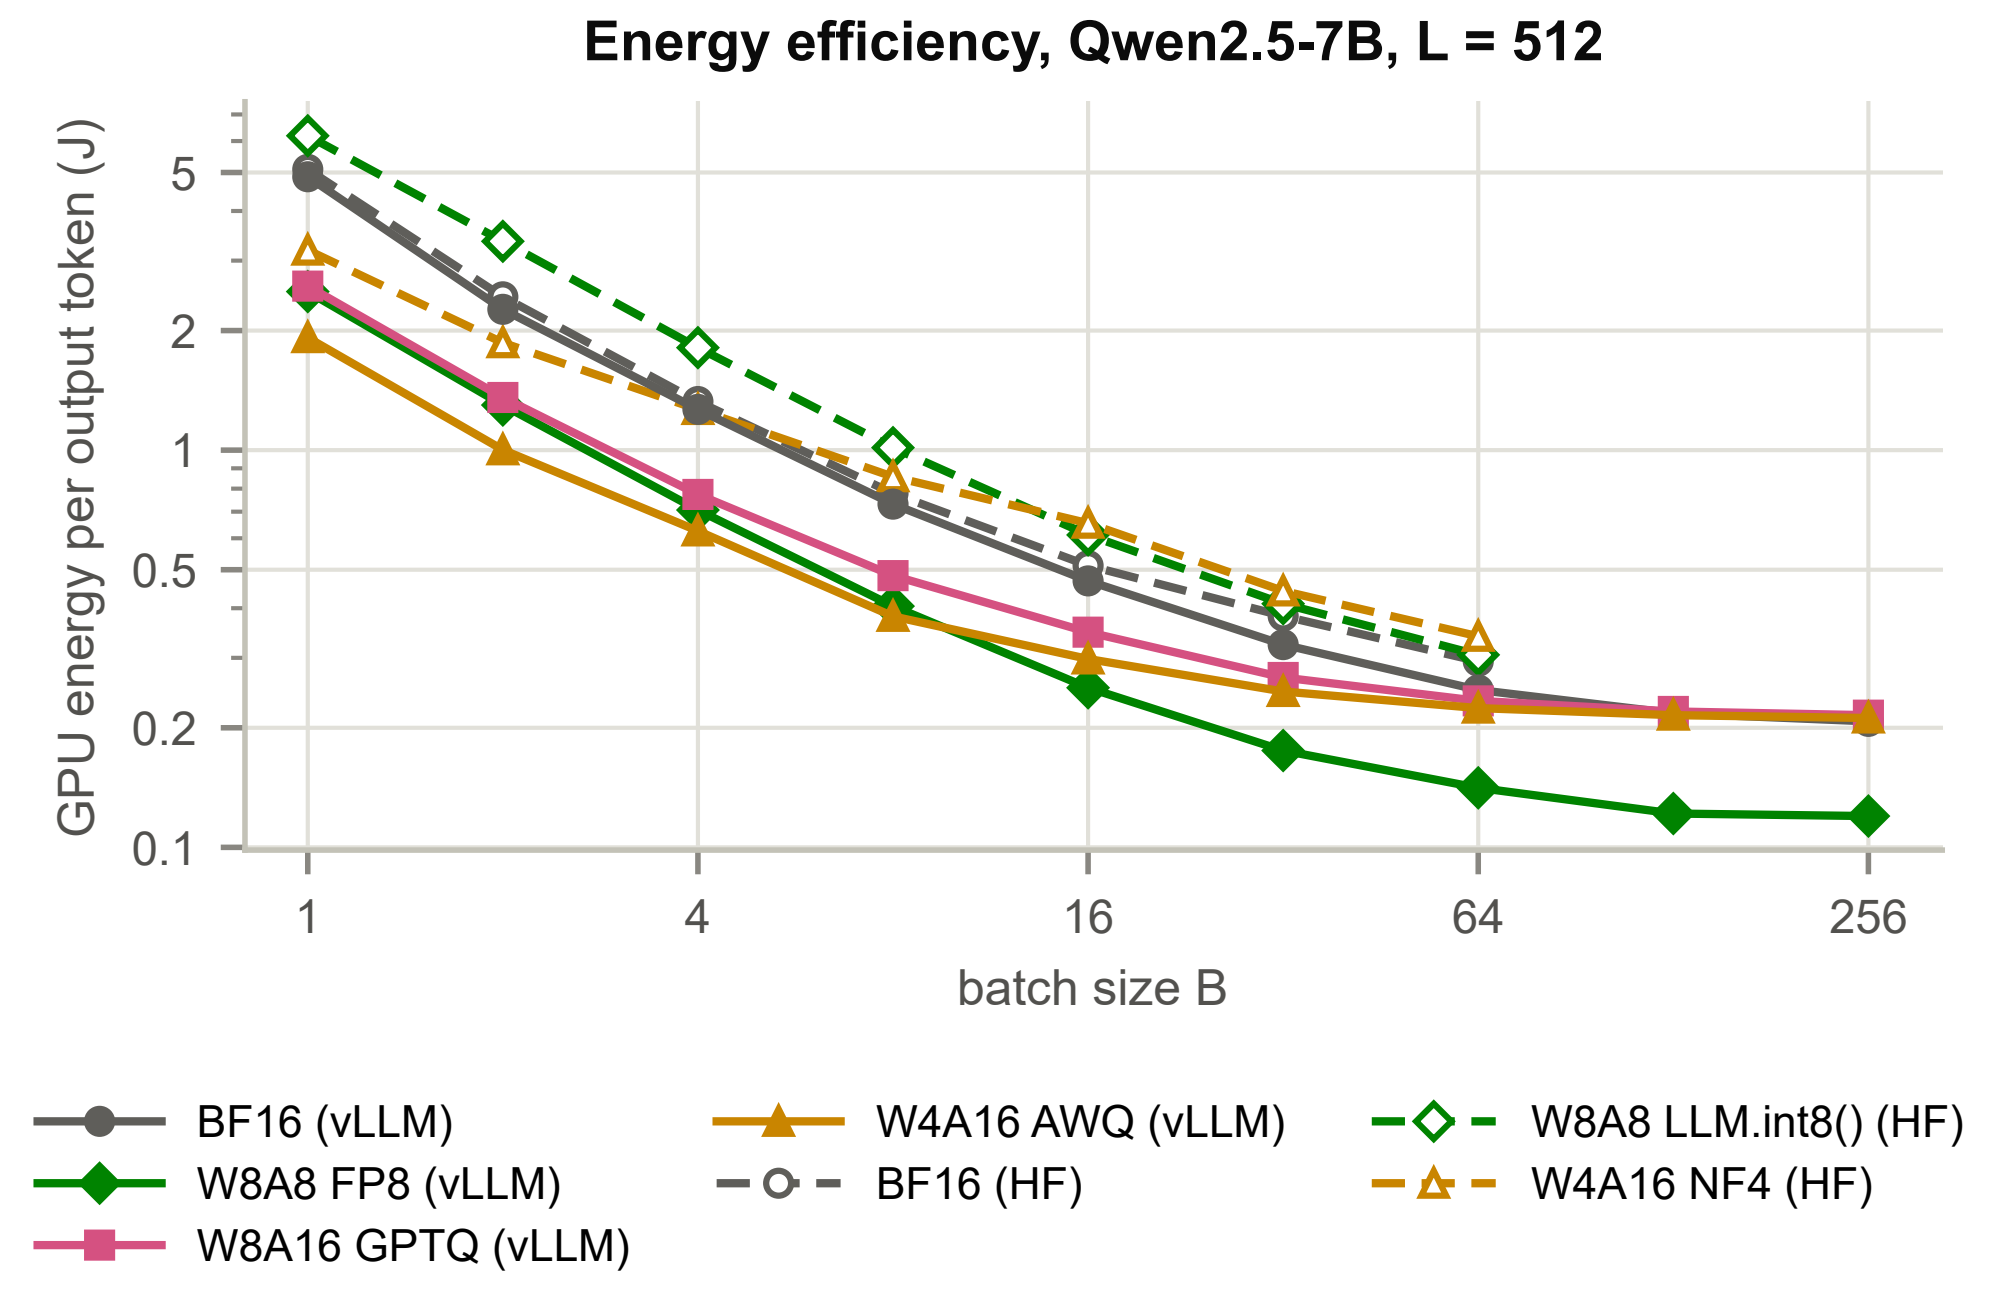
Energy efficiency, Qwen2.5-7B, L = 512
GPU energy per output token (J)
5
2
1
0.5
0.2
0.1
1
4
16
64
256
batch size B
BF16 (vLLM)
W8A8 FP8 (vLLM)
W8A16 GPTQ (vLLM)
W4A16 AWQ (vLLM)
BF16 (HF)
W8A8 LLM.int8() (HF)
W4A16 NF4 (HF)