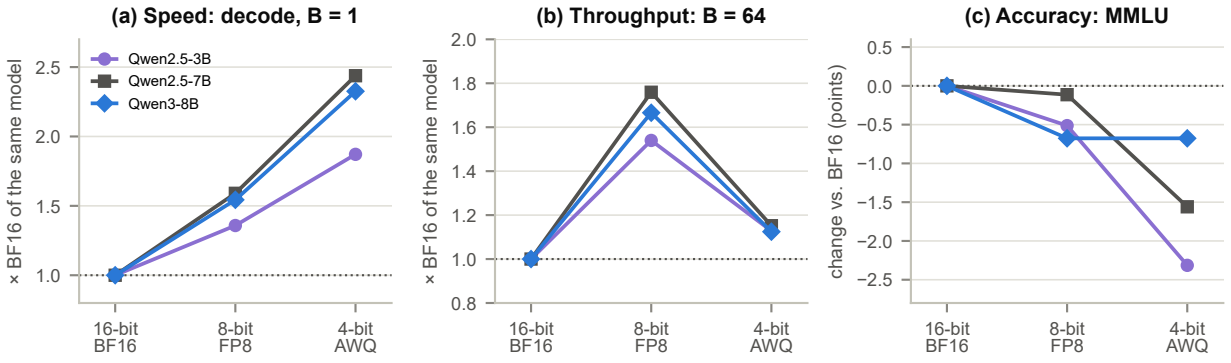
(a) Speed: decode, B = 1
(b) Throughput: B = 64
(c) Accuracy: MMLU
Qwen2.5-3B
Qwen2.5-7B
Qwen3-8B
× BF16 of the same model
change vs. BF16 (points)
2.5
2.0
1.5
1.0
0.5
0.0
−0.5
−1.0
−1.5
−2.0
−2.5
1.8
1.6
1.4
1.2
0.8
16-bit BF16
8-bit FP8
4-bit AWQ