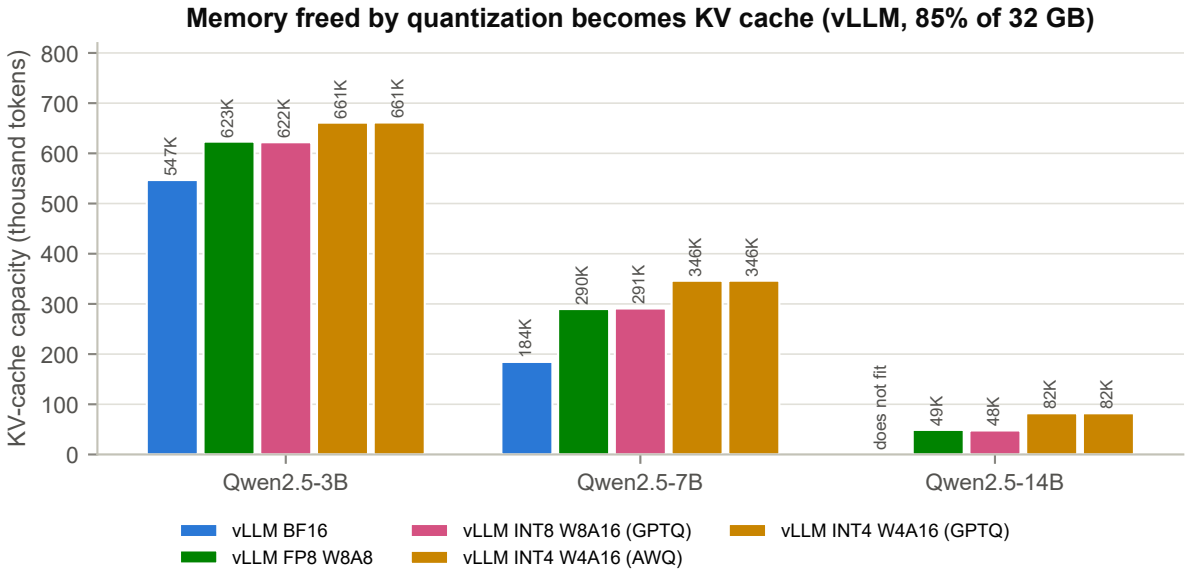
Memory freed by quantization becomes KV cache (vLLM, 85% of 32 GB)
KV-cache capacity (thousand tokens)
800
700
600
500
400
300
200
100
0
547K
623K
622K
661K
661K
184K
290K
291K
346K
346K
does not fit
49K
48K
82K
82K
Qwen2.5-3B
Qwen2.5-7B
Qwen2.5-14B
vLLM BF16
vLLM FP8 W8A8
vLLM INT8 W8A16 (GPTQ)
vLLM INT4 W4A16 (AWQ)
vLLM INT4 W4A16 (GPTQ)