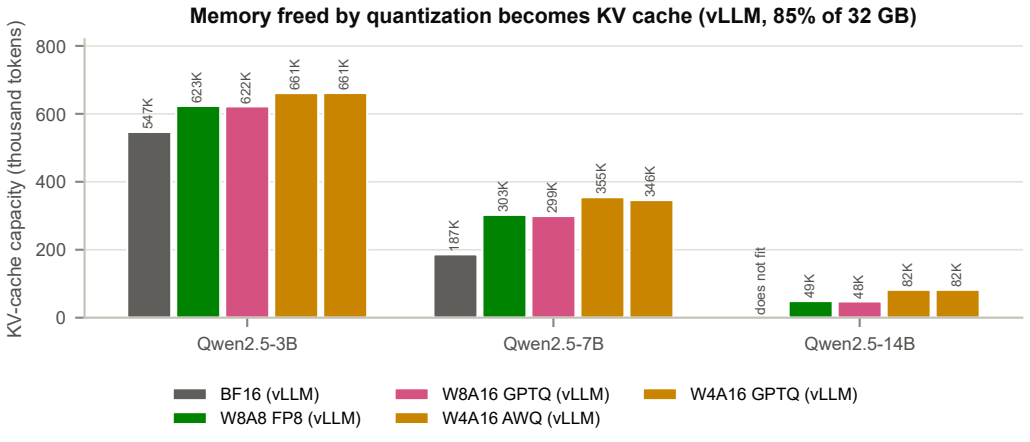
Memory freed by quantization becomes KV cache (vLLM, 85% of 32 GB)
KV-cache capacity (thousand tokens)
800
600
400
200
0
547K
623K
622K
661K
661K
187K
303K
299K
355K
346K
does not fit
49K
48K
82K
82K
Qwen2.5-3B
Qwen2.5-7B
Qwen2.5-14B
BF16 (vLLM)
W8A8 FP8 (vLLM)
W8A16 GPTQ (vLLM)
W4A16 AWQ (vLLM)
W4A16 GPTQ (vLLM)