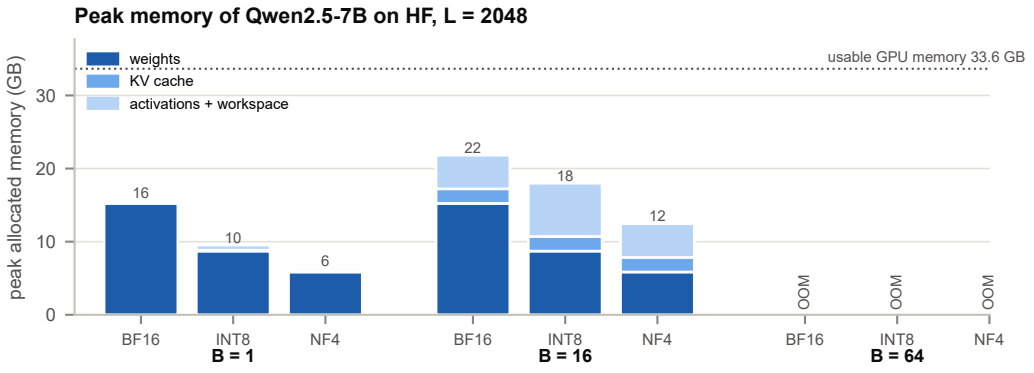
Peak memory of Qwen2.5-7B on HF, L = 2048
weights
KV cache
activations + workspace
usable GPU memory 33.6 GB
peak allocated memory (GB)
0
10
20
30
16
10
6
22
18
12
OOM
OOM
OOM
BF16
INT8
NF4
BF16
INT8
NF4
BF16
INT8
NF4
B = 1
B = 16
B = 64