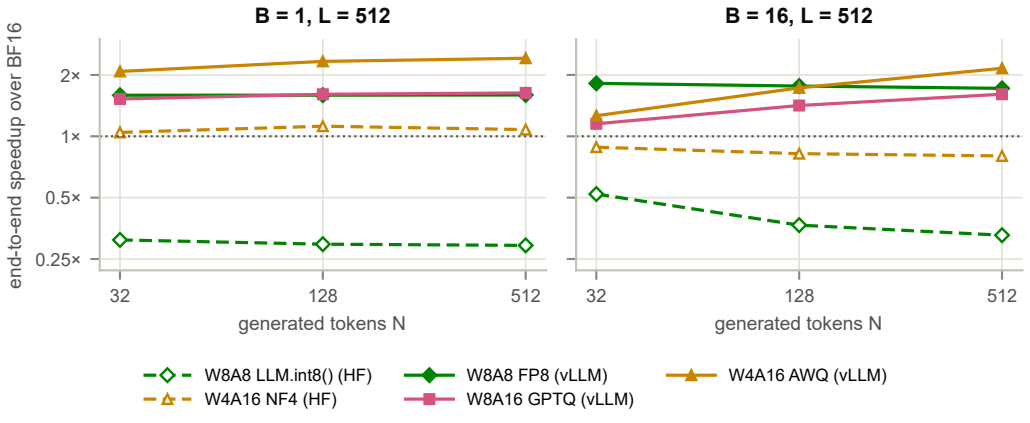
B = 1, L = 512
B = 16, L = 512
end-to-end speedup over BF16
2×
1×
0.5×
0.25×
32
128
512
generated tokens N
W8A8 LLM.int8() (HF)
W4A16 NF4 (HF)
W8A8 FP8 (vLLM)
W8A16 GPTQ (vLLM)
W4A16 AWQ (vLLM)