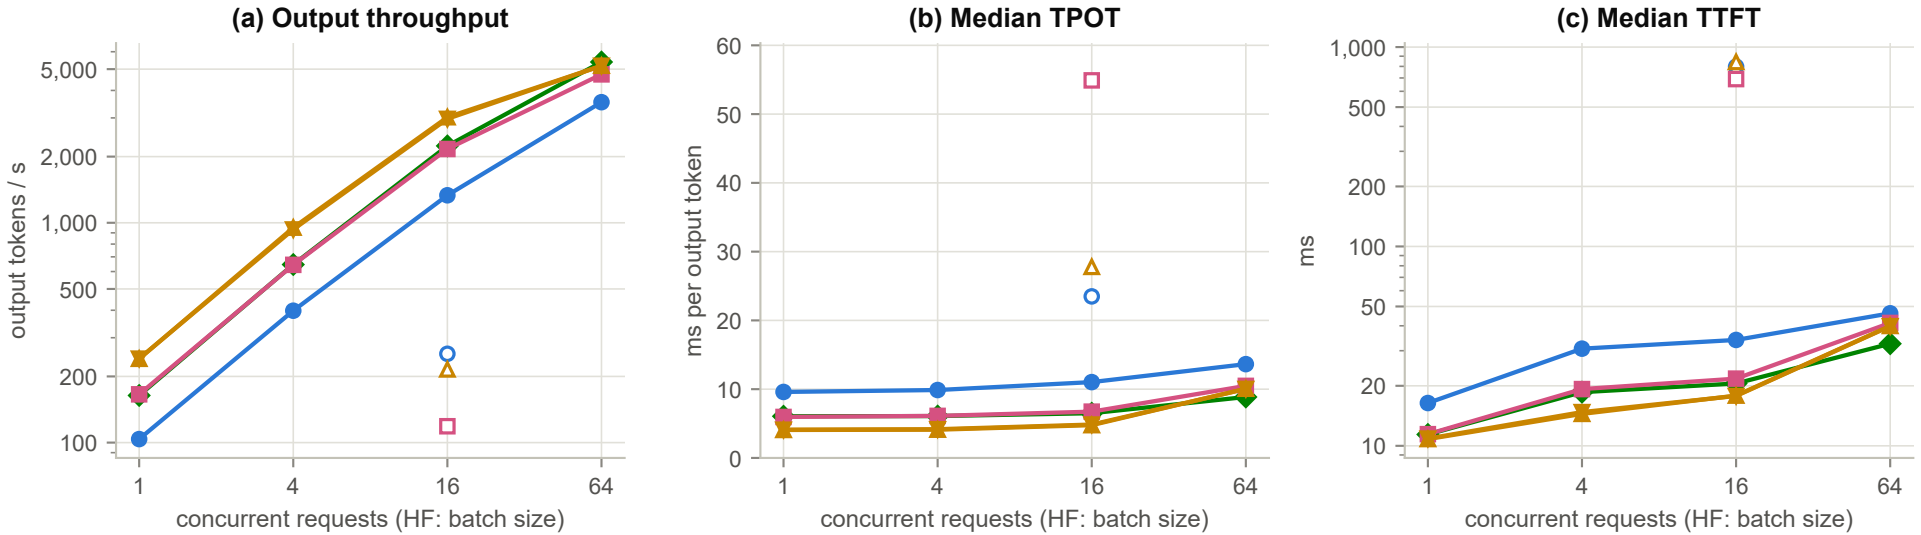
(a) Output throughput
output tokens / s
5,000
2,000
1,000
500
200
100
1
4
16
64
concurrent requests (HF: batch size)
(b) Median TPOT
ms per output token
60
50
40
30
20
10
0
1
4
16
64
concurrent requests (HF: batch size)
(c) Median TTFT
ms
1,000
500
200
100
50
20
10
1
4
16
64
concurrent requests (HF: batch size)
vLLM BF16
vLLM FP8 W8A8
vLLM INT8 W8A16 (GPTQ)
vLLM INT4 W4A16 (AWQ)
vLLM INT4 W4A16 (GPTQ)
HF BF16
HF INT8 (bnb LLM.int8)
HF NF4 (bnb)