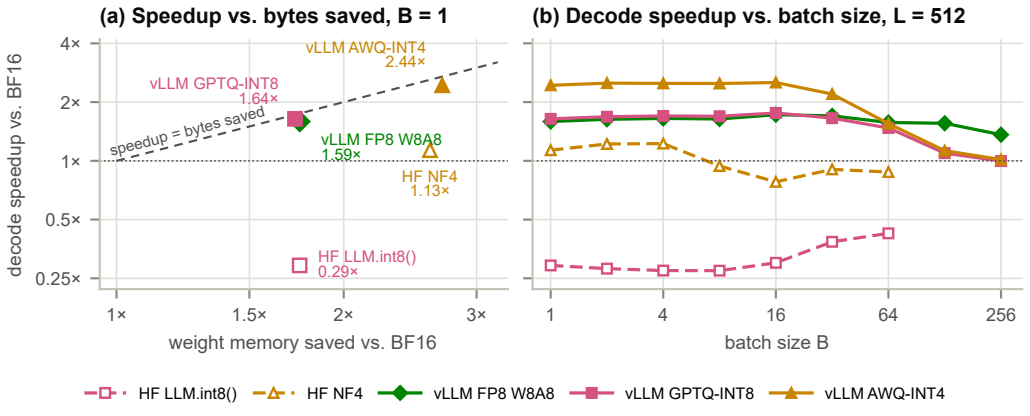
(a) Speedup vs. bytes saved, B = 1
(b) Decode speedup vs. batch size, L = 512
decode speedup vs. BF16
4×
2×
1×
0.5×
0.25×
vLLM AWQ-INT4
2.44×
vLLM GPTQ-INT8
1.64×
speedup = bytes saved
vLLM FP8 W8A8
1.59×
HF NF4
1.13×
HF LLM.int8()
0.29×
1×
1.5×
2×
3×
weight memory saved vs. BF16
1
4
16
64
256
batch size B
HF LLM.int8()
HF NF4
vLLM FP8 W8A8
vLLM GPTQ-INT8
vLLM AWQ-INT4